

104 Types of Flower Classification with TPUs

```
In [1]: import math, re, os
import tensorflow as tf
import numpy as np
from matplotlib import pyplot as plt
from sklearn.metrics import f1_score, precision_score, recall_score, confusion_matrix
print("Tensorflow version " + tf.__version__)
AUTO = tf.data.experimental.AUTOTUNE
```

D1022 21:38:22.909670386 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT tcp_frame_size_tuning
D1022 21:38:22.909693269 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT tcp_rcv_lowat
D1022 21:38:22.909696454 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT peer_state_based_framing
D1022 21:38:22.909698924 ON (default:ON)	15 config.cc:119]	gRPC EXPERIMENT flow_control_fixes
D1022 21:38:22.909701209 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT memory_pressure_controller
D1022 21:38:22.909705062 ize OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT unconstrained_max_quota_buffer_s
D1022 21:38:22.909707489 ON (default:ON)	15 config.cc:119]	gRPC EXPERIMENT new_hpack_huffman_decoder
D1022 21:38:22.909709786 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT event_engine_client
D1022 21:38:22.909712017 ON (default:ON)	15 config.cc:119]	gRPC EXPERIMENT monitoring_experiment
D1022 21:38:22.909714262 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT promise_based_client_call
D1022 21:38:22.909716429 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT free_large_allocator
D1022 21:38:22.909718590 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT promise_based_server_call
D1022 21:38:22.909720908 y OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT transport_supplies_client_latenc
D1022 21:38:22.909723016 OFF (default:OFF)	15 config.cc:119]	gRPC EXPERIMENT event_engine_listener
I1022 21:38:22.909945444	15 ev_epoll1_linux.cc:122]	grpc epoll fd: 62
D1022 21:38:22.909956923	15 ev_posix.cc:144]	Using polling engine: epoll1
D1022 21:38:22.909976772	15 dns_resolver_ares.cc:822]	Using ares dns resolver
D1022 21:38:22.922863174 rimental"	15 lb_policy_registry.cc:46]	registering LB policy factory for "priority_expe
D1022 21:38:22.922876805 tion_experimental"	15 lb_policy_registry.cc:46]	registering LB policy factory for "outlier_detec
D1022 21:38:22.922881066 et_experimental"	15 lb_policy_registry.cc:46]	registering LB policy factory for "weighted_targ
D1022 21:38:22.922883800	15 lb_policy_registry.cc:46]	registering LB policy factory for "pick_first"
D1022 21:38:22.922886599	15 lb_policy_registry.cc:46]	registering LB policy factory for "round_robin"
D1022 21:38:22.922889561 d_robin_experimental"	15 lb_policy_registry.cc:46]	registering LB policy factory for "weighted_roun
D1022 21:38:22.922897101 erimental"	15 lb_policy_registry.cc:46]	registering LB policy factory for "ring_hash_exp
D1022 21:38:22.922916472	15 lb_policy_registry.cc:46]	registering LB policy factory for "grpclb"
D1022 21:38:22.922946684	15 lb_policy_registry.cc:46]	registering LB policy factory for "rls_experimen

```

tal"
D1022 21:38:22.922962324      15 lb_policy_registry.cc:46]
anager_experimental"
D1022 21:38:22.922965838      15 lb_policy_registry.cc:46]
mpl_experimental"
D1022 21:38:22.922968832      15 lb_policy_registry.cc:46]
tal"
D1022 21:38:22.922975834      15 lb_policy_registry.cc:46]
esolver_experimental"
D1022 21:38:22.922978941      15 lb_policy_registry.cc:46]
host_experimental"
D1022 21:38:22.922981752      15 lb_policy_registry.cc:46]
ity_experimental"
D1022 21:38:22.922985932      15 certificate_provider_registry.cc:35] registering certificate provider factory for "fi
le_watcher"
I1022 21:38:22.926493876      15 socket_utils_common_posix.cc:408] Disabling AF_INET6 sockets because ::1 is not av
ailable.
I1022 21:38:22.940522295      240 socket_utils_common_posix.cc:337] TCP_USER_TIMEOUT is available. TCP_USER_TIMEOUT
will be used thereafter
E1022 21:38:22.947372221      240 oauth2_credentials.cc:236] oauth_fetch: UNKNOWN:C-ares status is not ARES_S
UCCESS qtype=A name=metadata.google.internal. is_balancer=0: Domain name not found {grpc_status:2, created_time:"2023-1
-0-22T21:38:22.947355967+00:00"}

```

Tensorflow version 2.12.0

TPU or GPU detection

In [2]: # NEW on TPU in TensorFlow 24: shorter cross-compatible TPU/GPU/multi-GPU/cluster-GPU detection code

```

try: # detect TPUs
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver.connect() # TPU detection
    strategy = tf.distribute.TPUStrategy(tpu)
except ValueError: # detect GPUs
    strategy = tf.distribute.MirroredStrategy() # for GPU or multi-GPU machines
    #strategy = tf.distribute.get_strategy() # default strategy that works on CPU and single GPU
    #strategy = tf.distribute.experimental.MultiWorkerMirroredStrategy() # for clusters of multi-GPU machines

print("Number of accelerators: ", strategy.num_replicas_in_sync)

```

```
INFO:tensorflow:Deallocate tpu buffers before initializing tpu system.  
INFO:tensorflow:Initializing the TPU system: local  
INFO:tensorflow:Finished initializing TPU system.  
INFO:tensorflow:Found TPU system:  
INFO:tensorflow:*** Num TPU Cores: 8  
INFO:tensorflow:*** Num TPU Workers: 1  
INFO:tensorflow:*** Num TPU Cores Per Worker: 8  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:CPU:0, CPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:0, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:1, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:2, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:3, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:4, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:5, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:6, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU:7, TPU, 0, 0)  
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:TPU_SYSTEM:0, TPU_SYSTEM, 0, 0)  
Number of accelerators: 8
```

Configuration

```

'barberton daisy',  'daffodil',
'petunia',          'wild pansy',
'cautleya spicata', 'japanese anemone',
'gazania',          'azalea',
'frangipani',       'clematis',
'hippeastrum ',     'bee balm',
'trumpet creeper',  'blackberry lily',
'sword lily',        'poinsettia',
'primula',          'sunflower',
'black-eyed susan', 'silverbush',
'water lily',        'rose',
'hibiscus',          'columbine',
'pink quill',        'foxglove',
'common tulip',      'wild rose']

```

Visualization utilities

```

In [5]: # numpy and matplotlib defaults
np.set_printoptions(threshold=15, linewidth=80)

def batch_to_numpy_images_and_labels(data):
    images, labels = data
    numpy_images = images.numpy()
    numpy_labels = labels.numpy()
    if numpy_labels.dtype == object: # binary string in this case, these are image ID strings
        numpy_labels = [None for _ in enumerate(numpy_images)]
    # If no labels, only image IDs, return None for labels (this is the case for test data)
    return numpy_images, numpy_labels

def title_from_label_and_target(label, correct_label):
    if correct_label is None:
        return CLASSES[label], True
    correct = (label == correct_label)
    return "{} [{}{}{}]".format(CLASSES[label], 'OK' if correct else 'NO', u"\u2192" if not correct else '',
                                CLASSES[correct_label] if not correct else ''), correct

def display_one_flower(image, title, subplot, red=False, titlesize=16):
    plt.subplot(*subplot)
    plt.axis('off')
    plt.imshow(image)
    if len(title) > 0:
        plt.title(title, fontsize=int(titlesize) if not red else int(titlesize/1.2), color='red' if red else 'black', fontweight='bold')
    return subplot[0], subplot[1], subplot[2]+1

def display_batch_of_images(databatch, predictions=None):
    """This will work with:
    display_batch_of_images(images)
    display_batch_of_images(images, predictions)
    display_batch_of_images((images, labels))
    display_batch_of_images((images, labels), predictions)
    """
    if isinstance(databatch, tuple):
        images, labels = databatch
    else:
        images, labels = databatch.images, databatch.labels if hasattr(databatch, 'labels') else None
    n_images, n_rows, n_cols = images.shape[0], images.shape[1], images.shape[2]
    if labels is None:
        labels = [None]*n_images
    if predictions is None:
        predictions = [None]*n_images
    # Every row is a subplot
    subplots = [display_one_flower(images[i], title=(labels[i], predictions[i]), subplot=(i//n_cols, i%n_cols, i+1)) for i in range(n_images)]
    return subplots

```

```

"""
# data
images, labels = batch_to_numpy_images_and_labels(databatch)
if labels is None:
    labels = [None for _ in enumerate(images)]

# auto-squaring: this will drop data that does not fit into square or square-ish rectangle
rows = int(math.sqrt(len(images)))
cols = len(images)//rows

# size and spacing
FIGSIZE = 13.0
SPACING = 0.1
subplot=(rows,cols,1)
if rows < cols:
    plt.figure(figsize=(FIGSIZE,FIGSIZE/cols*rows))
else:
    plt.figure(figsize=(FIGSIZE/rows*cols,FIGSIZE))

# display
for i, (image, label) in enumerate(zip(images[:rows*cols], labels[:rows*cols])):
    title = '' if label is None else CLASSES[label]
    correct = True
    if predictions is not None:
        title, correct = title_from_label_and_target(predictions[i], label)
    dynamic_titlesize = FIGSIZE*SPACING/max(rows,cols)*40+3 # magic formula tested to work from 1x1 to 10x10 images
    subplot = display_one_flower(image, title, subplot, not correct, titlesize=dynamic_titlesize)

#Layout
plt.tight_layout()
if label is None and predictions is None:
    plt.subplots_adjust(wspace=0, hspace=0)
else:
    plt.subplots_adjust(wspace=SPACING, hspace=SPACING)
plt.show()

def display_confusion_matrix(cmat, score, precision, recall):
    plt.figure(figsize=(15,15))
    ax = plt.gca()
    ax.matshow(cmat, cmap='Reds')
    ax.set_xticks(range(len(CLASSES)))
    ax.set_xticklabels(CLASSES, fontdict={'fontsize': 7})
    plt.setp(ax.get_xticklabels(), rotation=45, ha="left", rotation_mode="anchor")
    ax.set_yticks(range(len(CLASSES)))
    ax.set_yticklabels(CLASSES, fontdict={'fontsize': 7})

```

```

plt.setp(ax.get_yticklabels(), rotation=45, ha="right", rotation_mode="anchor")
titlestring = ""
if score is not None:
    titlestring += 'f1 = {:.3f} '.format(score)
if precision is not None:
    titlestring += '\nprecision = {:.3f} '.format(precision)
if recall is not None:
    titlestring += '\nrecall = {:.3f} '.format(recall)
if len(titlestring) > 0:
    ax.text(101, 1, titlestring, fontdict={'fontsize': 18, 'horizontalalignment':'right', 'verticalalignment':'top'})
plt.show()

def display_training_curves(training, validation, title, subplot):
    if subplot%10==1: # set up the subplots on the first call
        plt.subplots(figsize=(10,10), facecolor="#F0F0F0")
        plt.tight_layout()
    ax = plt.subplot(subplot)
    ax.set_facecolor('#F8F8F8')
    ax.plot(training)
    ax.plot(validation)
    ax.set_title('model ' + title)
    ax.set_ylabel(title)
    #ax.set_ylim(0.28,1.05)
    ax.set_xlabel('epoch')
    ax.legend(['train', 'valid.'])

```

Datasets

In [6]:

```

def decode_image(image_data):
    image = tf.image.decode_jpeg(image_data, channels=3) # image format uint8 [0,255]
    image = tf.reshape(image, [*IMAGE_SIZE, 3]) # explicit size needed for TPU
    return image

def read_labeled_tfrecord(example):
    LABELED_TFREC_FORMAT = {
        "image": tf.io.FixedLenFeature([], tf.string), # tf.string means bytestring
        "class": tf.io.FixedLenFeature([], tf.int64), # shape [] means single element
    }
    example = tf.io.parse_single_example(example, LABELED_TFREC_FORMAT)
    image = decode_image(example['image'])
    label = tf.cast(example['class'], tf.int32)
    return image, label # returns a dataset of (image, label) pairs

```

```

def read_unlabeled_tfrecord(example):
    UNLABELED_TFREC_FORMAT = {
        "image": tf.io.FixedLenFeature([], tf.string), # tf.string means bytestring
        "id": tf.io.FixedLenFeature([], tf.string), # shape [] means single element
        # class is missing, this competition's challenge is to predict flower classes for the test dataset
    }
    example = tf.io.parse_single_example(example, UNLABELED_TFREC_FORMAT)
    image = decode_image(example['image'])
    idnum = example['id']
    return image, idnum # returns a dataset of image(s)

def load_dataset(filenames, labeled=True, ordered=False):
    # Read from TFRecords. For optimal performance, reading from multiple files at once and
    # disregarding data order. Order does not matter since we will be shuffling the data anyway.

    ignore_order = tf.data.Options()
    if not ordered:
        ignore_order.experimental_deterministic = False # disable order, increase speed

    dataset = tf.data.TFRecordDataset(filenames, num_parallel_reads=AUTO) # automatically interleaves reads from multiple
    dataset = dataset.with_options(ignore_order) # uses data as soon as it streams in, rather than in its original order
    dataset = dataset.map(read_labeled_tfrecord if labeled else read_unlabeled_tfrecord, num_parallel_calls=AUTO)
    # returns a dataset of (image, label) pairs if Labeled=True or (image, id) pairs if Labeled=False
    return dataset

def data_augment(image, label):
    # data augmentation. Thanks to the dataset.prefetch(AUTO) statement in the next function (below),
    # this happens essentially for free on TPU. Data pipeline code is executed on the "CPU" part
    # of the TPU while the TPU itself is computing gradients.
    image = tf.image.random_flip_left_right(image)
    #image = tf.image.random_saturation(image, 0, 2)
    return image, label

def get_training_dataset():
    dataset = load_dataset(TRAINING_FILENAMES, labeled=True)
    dataset = dataset.map(data_augment, num_parallel_calls=AUTO)
    dataset = dataset.repeat() # the training dataset must repeat for several epochs
    dataset = dataset.shuffle(2048)
    dataset = dataset.batch(BATCH_SIZE)
    dataset = dataset.prefetch(AUTO) # prefetch next batch while training (autotune prefetch buffer size)
    return dataset

def get_validation_dataset(ordered=False):
    dataset = load_dataset(VALIDATION_FILENAMES, labeled=True, ordered=ordered)
    dataset = dataset.batch(BATCH_SIZE)

```

```

dataset = dataset.cache()
dataset = dataset.prefetch(AUTO) # prefetch next batch while training (autotune prefetch buffer size)
return dataset

def get_test_dataset(ordered=False):
    dataset = load_dataset(TEST_Filenames, labeled=False, ordered=ordered)
    dataset = dataset.batch(BATCH_SIZE)
    dataset = dataset.prefetch(AUTO) # prefetch next batch while training (autotune prefetch buffer size)
    return dataset

def count_data_items(filenames):
    # the number of data items is written in the name of the .tfrec files, i.e. flowers00-230.tfrec = 230 data items
    n = [int(re.compile(r"-([0-9]*)\.").search(filename).group(1)) for filename in filenames]
    return np.sum(n)

NUM_TRAINING_IMAGES = count_data_items(TRAINING_Filenames)
NUM_VALIDATION_IMAGES = count_data_items(VALIDATION_Filenames)
NUM_TEST_IMAGES = count_data_items(TEST_Filenames)
STEPS_PER_EPOCH = NUM_TRAINING_IMAGES // BATCH_SIZE
VALIDATION_STEPS = -(-NUM_VALIDATION_IMAGES // BATCH_SIZE) # The "-(-//)" trick rounds up instead of down :-
TEST_STEPS = -(-NUM_TEST_IMAGES // BATCH_SIZE) # The "-(-//)" trick rounds up instead of down :-
print('Dataset: {} training images, {} validation images, {} unlabeled test images'.format(NUM_TRAINING_IMAGES, NUM_VALIDATION_IMAGES, NUM_TEST_IMAGES))

```

Dataset: 12753 training images, 3712 validation images, 7382 unlabeled test images

Dataset visualizations

```

In [7]: # data dump
print("Training data shapes:")
for image, label in get_training_dataset().take(3):
    print(image.numpy().shape, label.numpy().shape)
print("Training data label examples:", label.numpy())
print("Validation data shapes:")
for image, label in get_validation_dataset().take(3):
    print(image.numpy().shape, label.numpy().shape)
print("Validation data label examples:", label.numpy())
print("Test data shapes:")
for image, idnum in get_test_dataset().take(3):
    print(image.numpy().shape, idnum.numpy().shape)
print("Test data IDs:", idnum.numpy().astype('U')) # U=unicode string

```

```
Training data shapes:  
(128, 512, 512, 3) (128,)  
(128, 512, 512, 3) (128,)  
(128, 512, 512, 3) (128,)  
Training data label examples: [73 26 62 ... 88 50 46]  
Validation data shapes:  
(128, 512, 512, 3) (128,)  
(128, 512, 512, 3) (128,)  
(128, 512, 512, 3) (128,)  
Validation data label examples: [103 53 49 ... 73 95 88]  
Test data shapes:  
(128, 512, 512, 3) (128,)  
(128, 512, 512, 3) (128,)  
(128, 512, 512, 3) (128,)  
Test data IDs: ['b51c9863e' '92c33e670' '8756842ba' ... '2b782bb3f' '23253eaed' 'd25115245']
```

```
In [8]: # Peek at training data  
training_dataset = get_training_dataset()  
training_dataset = training_dataset.unbatch().batch(20)  
train_batch = iter(training_dataset)
```

```
In [9]: # run this cell again for next set of images  
display_batch_of_images(next(train_batch))
```

yellow iris



tree poppy



hard-leaved pocket orchid



daisy



azalea



common dandelion



wild geranium



primula



common tulip



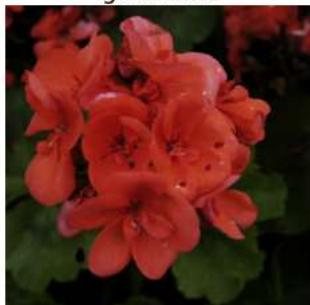
spear thistle



sweet william



geranium



petunia



sunflower



buttercup



sword lily



wild rose



mexican petunia



wild geranium



iris



```
In [10]: # peer at test data
test_dataset = get_test_dataset()
test_dataset = test_dataset.unbatch().batch(20)
test_batch = iter(test_dataset)
```

```
In [11]: # run this cell again for next set of images
display_batch_of_images(next(test_batch))
```



Model

```
In [12]: with strategy.scope():
    #img_adjust_layer = tf.keras.layers.Lambda(lambda data: tf.keras.applications.xception.preprocess_input(tf.cast(data, tf.float32)))
    #pretrained_model = tf.keras.applications.Xception(weights='imagenet', include_top=False)

    img_adjust_layer = tf.keras.layers.Lambda(lambda data: tf.keras.applications.vgg16.preprocess_input(tf.cast(data, tf.float32)))
    pretrained_model = tf.keras.applications.VGG16(weights='imagenet', include_top=False)

    pretrained_model.trainable = False # False = transfer Learning, True = fine-tuning

    model = tf.keras.Sequential([
        img_adjust_layer,
        pretrained_model,
        tf.keras.layers.GlobalAveragePooling2D(),
        tf.keras.layers.Dense(len(CLASSES), activation='softmax')
    ])

    model.compile(
        optimizer='adam',
        loss = 'sparse_categorical_crossentropy',
        metrics=['sparse_categorical_accuracy'],
        # NEW on TPU in TensorFlow 2.4: sending multiple batches to the TPU at once saves communications
        # overheads and allows the XLA compiler to unroll the loop on TPU and optimize hardware utilization.
        steps_per_execution=16
    )
    model.summary()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_\
tf_kernels_notop.h5
58889256/58889256 [=====] - 0s 0us/step
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====		
lambda (Lambda)	(None, 512, 512, 3)	0
vgg16 (Functional)	(None, None, None, 512)	14714688
global_average_pooling2d (GloballyAveragePooling2D)	(None, 512)	0
dense (Dense)	(None, 104)	53352
=====		
Total params:	14,768,040	
Trainable params:	53,352	
Non-trainable params:	14,714,688	

Training

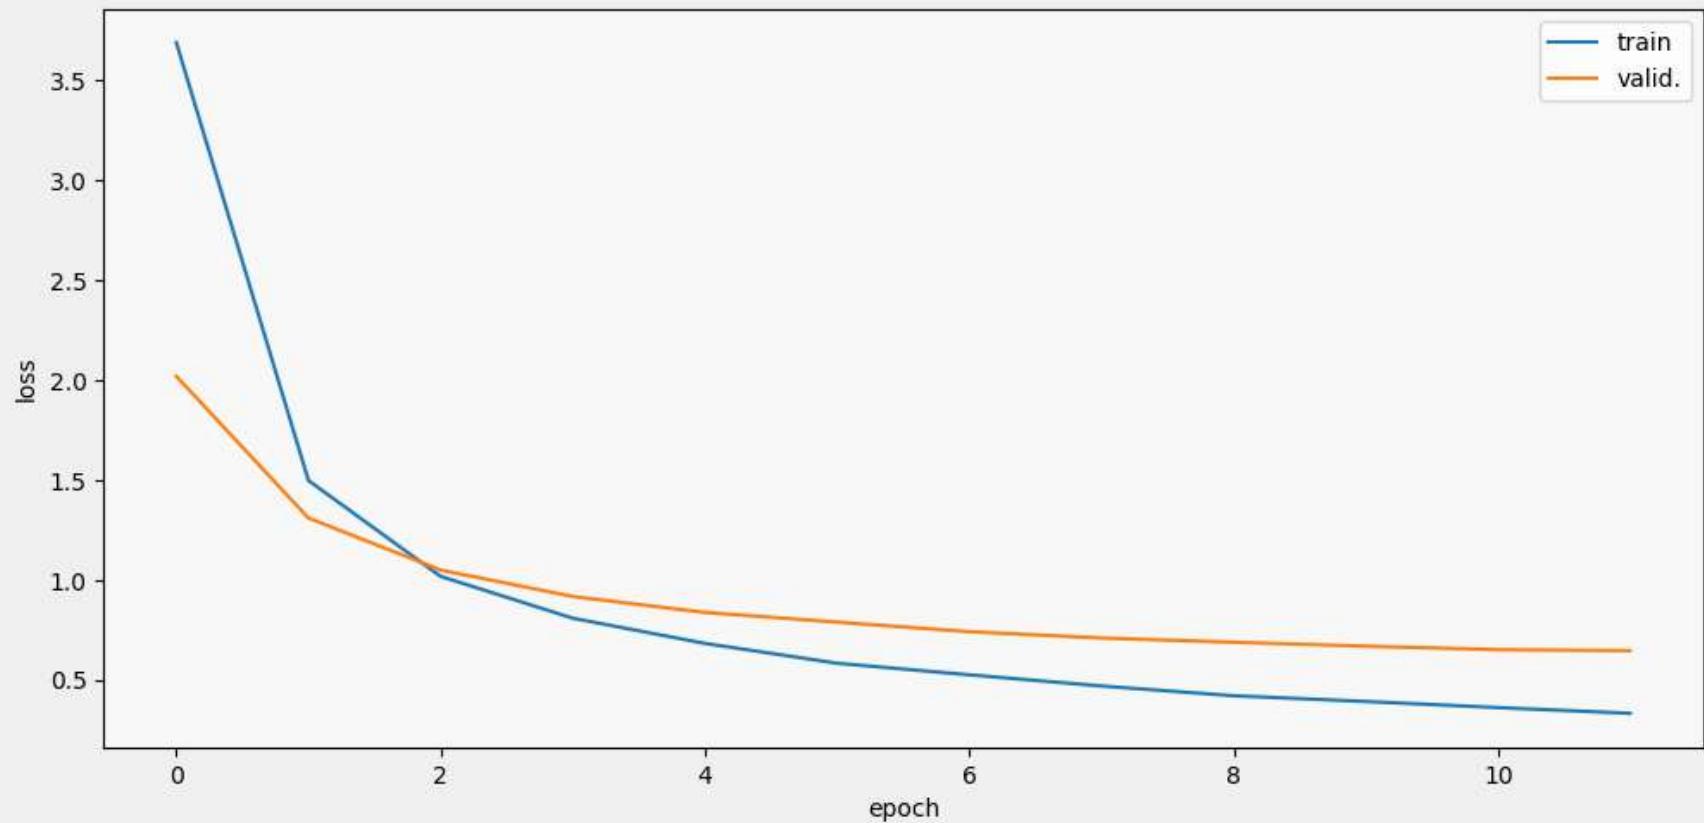
```
In [13]: history = model.fit(get_training_dataset(), steps_per_epoch=STEPS_PER_EPOCH, epochs=EPOCHS,
                           validation_data=get_validation_dataset(), validation_steps=VALIDATION_STEPS)
```

```
Epoch 1/12
99/99 [=====] - 30s 303ms/step - loss: 3.6911 - sparse_categorical_accuracy: 0.2514 - val_loss: 2.0196 - val_sparse_categorical_accuracy: 0.5116
Epoch 2/12
99/99 [=====] - 12s 122ms/step - loss: 1.4973 - sparse_categorical_accuracy: 0.6289 - val_loss: 1.3093 - val_sparse_categorical_accuracy: 0.6813
Epoch 3/12
99/99 [=====] - 12s 122ms/step - loss: 1.0170 - sparse_categorical_accuracy: 0.7504 - val_loss: 1.0488 - val_sparse_categorical_accuracy: 0.7452
Epoch 4/12
99/99 [=====] - 12s 122ms/step - loss: 0.8074 - sparse_categorical_accuracy: 0.8061 - val_loss: 0.9162 - val_sparse_categorical_accuracy: 0.7772
Epoch 5/12
99/99 [=====] - 12s 124ms/step - loss: 0.6813 - sparse_categorical_accuracy: 0.8363 - val_loss: 0.8364 - val_sparse_categorical_accuracy: 0.7939
Epoch 6/12
99/99 [=====] - 12s 122ms/step - loss: 0.5818 - sparse_categorical_accuracy: 0.8607 - val_loss: 0.7878 - val_sparse_categorical_accuracy: 0.8074
Epoch 7/12
99/99 [=====] - 12s 122ms/step - loss: 0.5239 - sparse_categorical_accuracy: 0.8777 - val_loss: 0.7399 - val_sparse_categorical_accuracy: 0.8195
Epoch 8/12
99/99 [=====] - 12s 122ms/step - loss: 0.4678 - sparse_categorical_accuracy: 0.8904 - val_loss: 0.7082 - val_sparse_categorical_accuracy: 0.8265
Epoch 9/12
99/99 [=====] - 12s 122ms/step - loss: 0.4190 - sparse_categorical_accuracy: 0.9009 - val_loss: 0.6872 - val_sparse_categorical_accuracy: 0.8330
Epoch 10/12
99/99 [=====] - 12s 122ms/step - loss: 0.3903 - sparse_categorical_accuracy: 0.9101 - val_loss: 0.6672 - val_sparse_categorical_accuracy: 0.8397
Epoch 11/12
99/99 [=====] - 12s 122ms/step - loss: 0.3597 - sparse_categorical_accuracy: 0.9176 - val_loss: 0.6498 - val_sparse_categorical_accuracy: 0.8408
Epoch 12/12
99/99 [=====] - 12s 122ms/step - loss: 0.3311 - sparse_categorical_accuracy: 0.9237 - val_loss: 0.6450 - val_sparse_categorical_accuracy: 0.8416
```

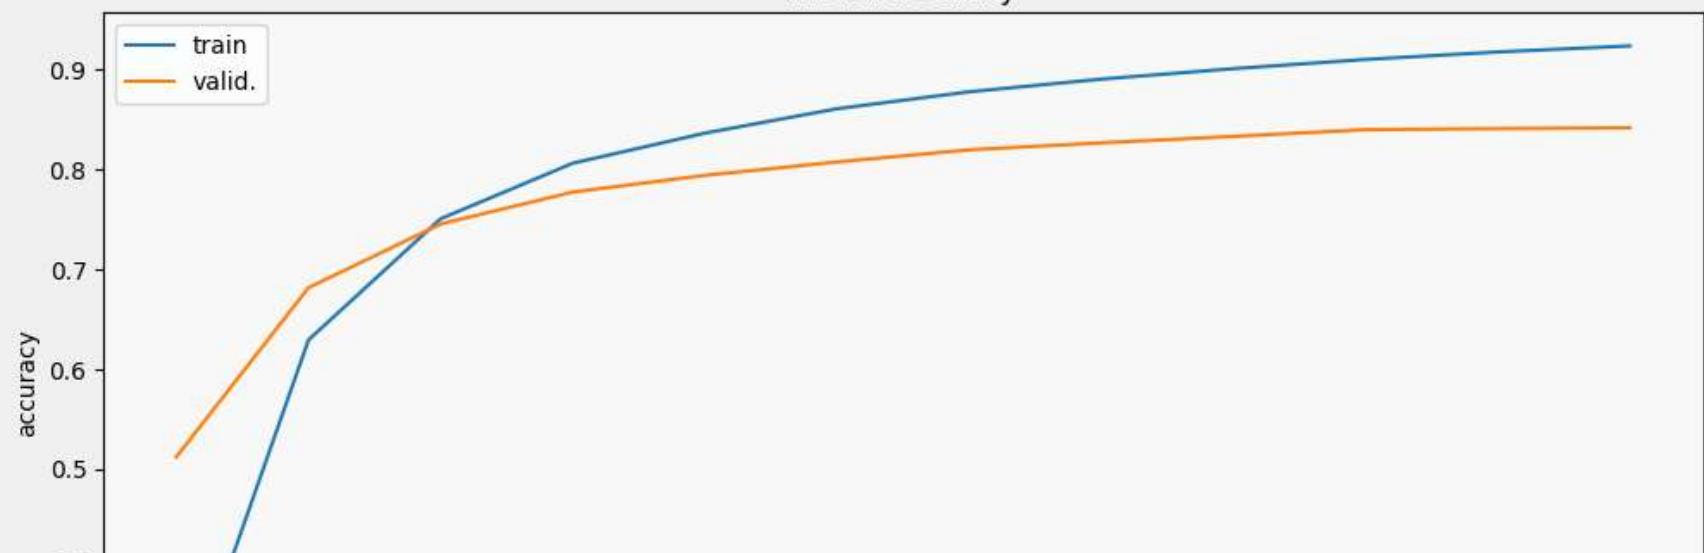
```
In [14]: display_training_curves(history.history['loss'], history.history['val_loss'], 'loss', 211)
display_training_curves(history.history['sparse_categorical_accuracy'], history.history['val_sparse_categorical_accuracy'], 'accuracy', 211)
```

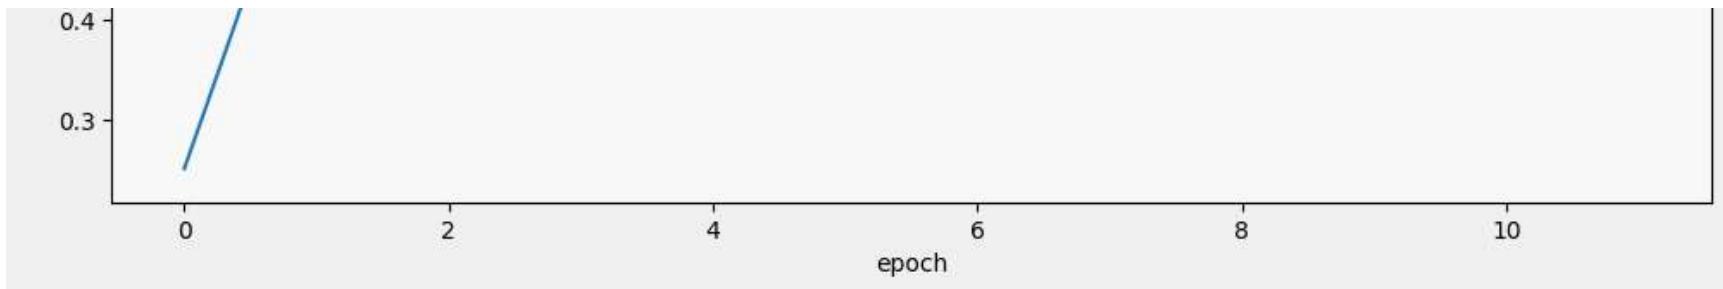
```
/tmp/ipykernel_15/3094237293.py:95: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.
  ax = plt.subplot(subplot)
```

model loss



model accuracy





Confusion matrix

```
In [15]: cmdataset = get_validation_dataset(ordered=True) # since we are splitting the dataset and iterating separately on images
images_ds = cmdataset.map(lambda image, label: image)
labels_ds = cmdataset.map(lambda image, label: label).unbatch()
cm_correct_labels = next(iter(labels_ds.batch(NUM_VALIDATION_IMAGES))).numpy() # get everything as one batch
cm_probabilities = model.predict(images_ds, steps=VALIDATION_STEPS)
cm_predictions = np.argmax(cm_probabilities, axis=-1)
print("Correct labels: ", cm_correct_labels.shape, cm_correct_labels)
print("Predicted labels: ", cm_predictions.shape, cm_predictions)
```

2023-10-22 21:41:40.661302: E tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] model_pruner failed: INVALID_ARGUMENT: Graph does not contain terminal node AssignAddVariableOp.

2023-10-22 21:41:40.855424: E tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] model_pruner failed: INVALID_ARGUMENT: Graph does not contain terminal node AssignAddVariableOp.

29/29 [=====] - 17s 593ms/step

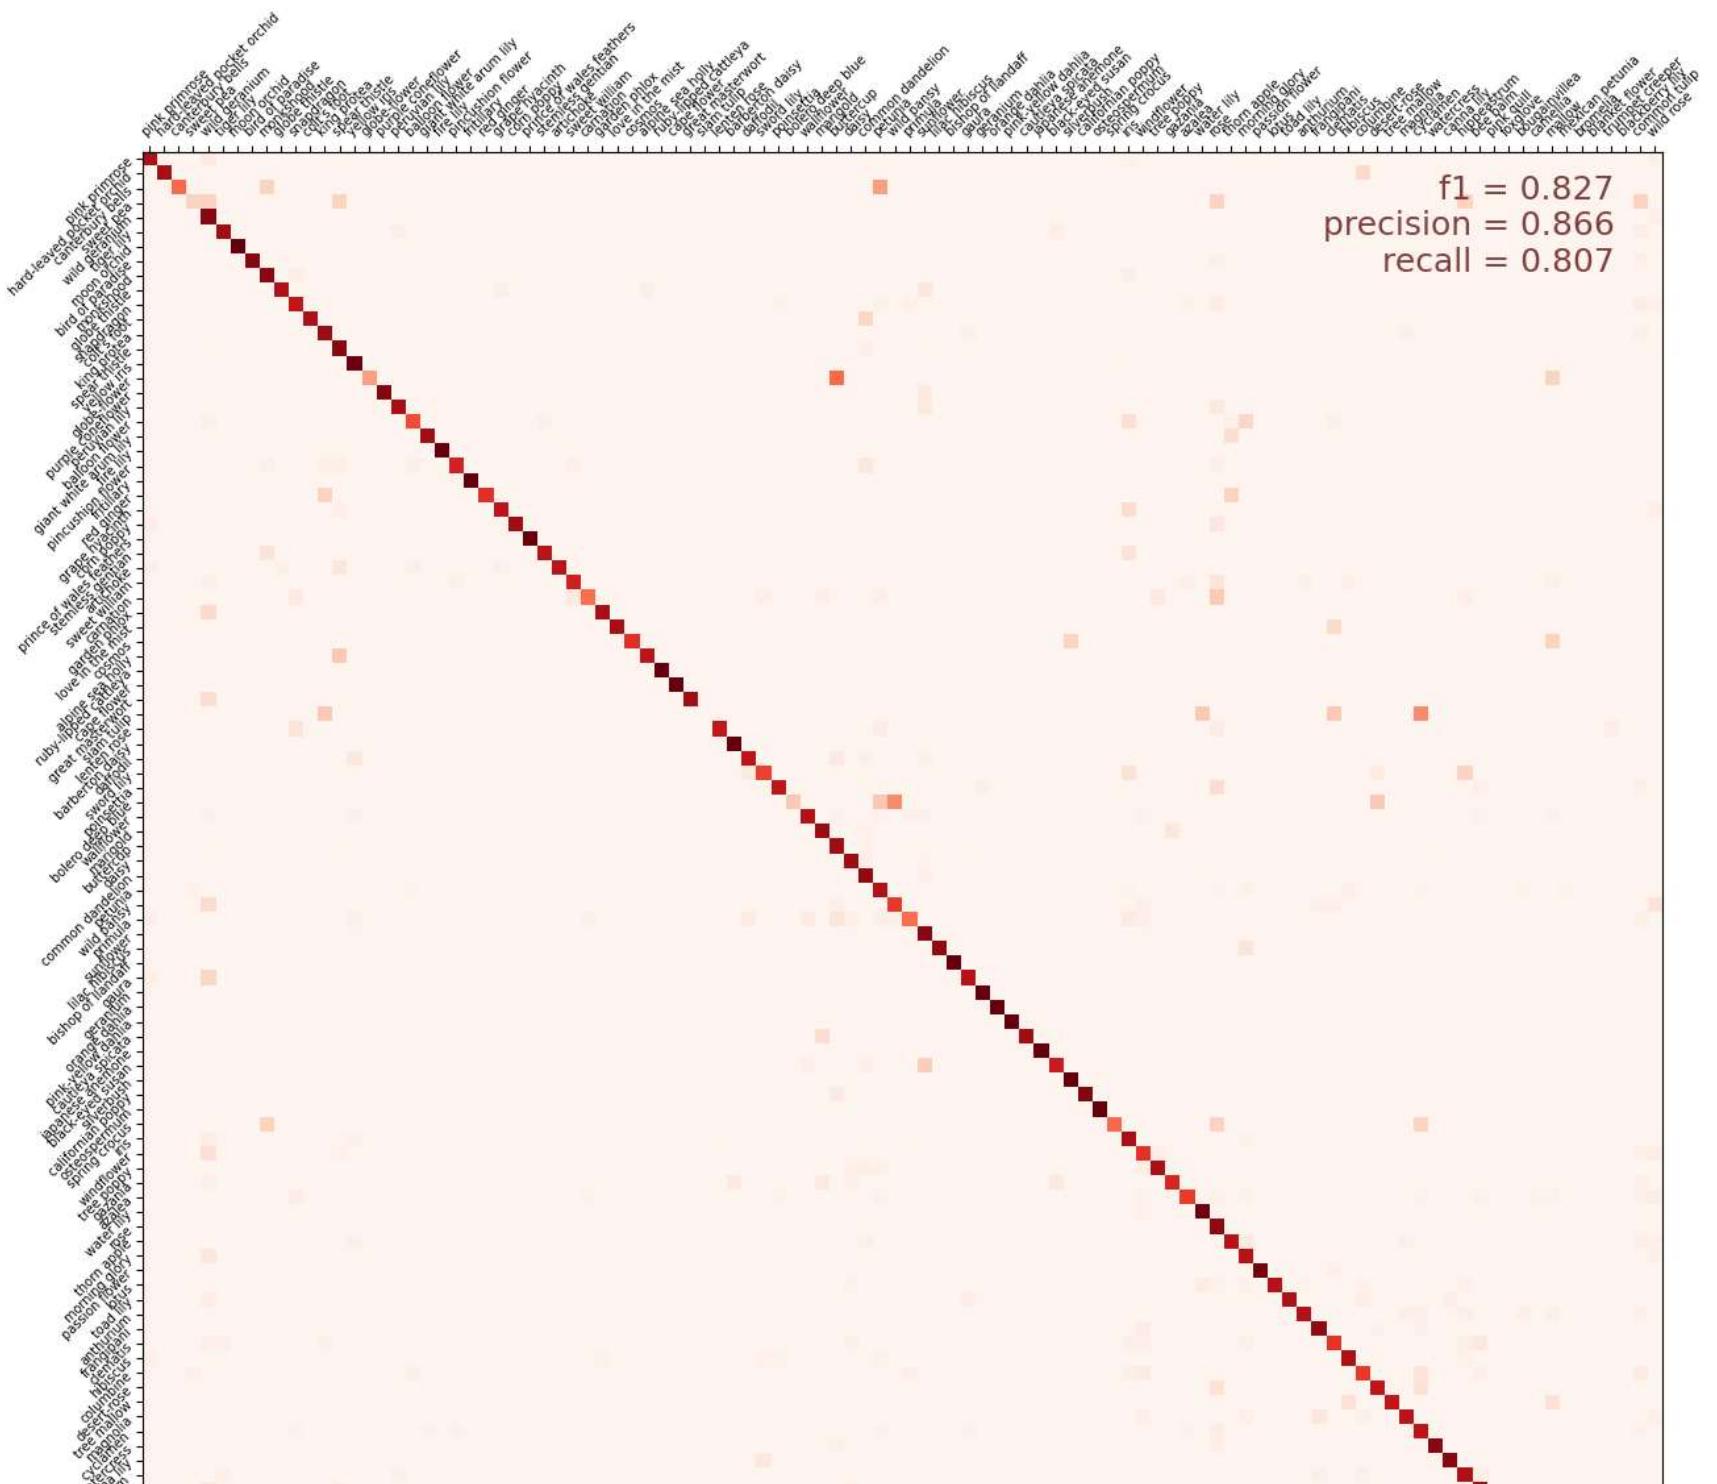
Correct labels: (3712,) [74 82 62 ... 67 50 53]

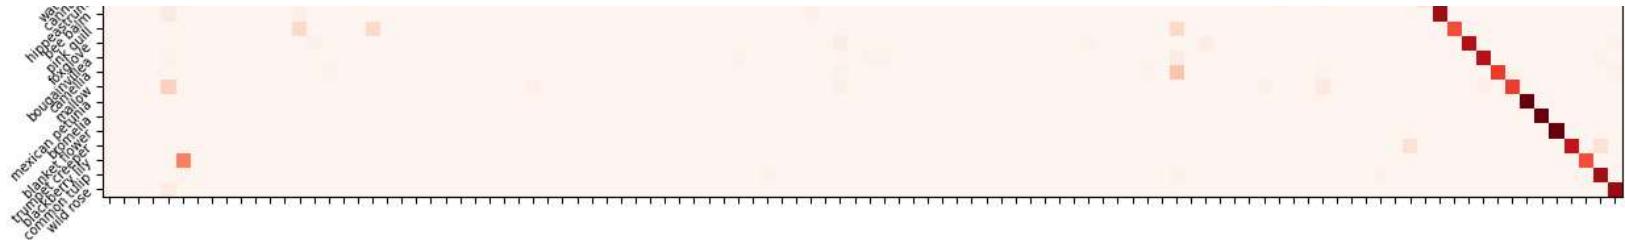
Predicted labels: (3712,) [74 82 53 ... 67 50 53]

```
In [16]: cmat = confusion_matrix(cm_correct_labels, cm_predictions, labels=range(len(CLASSES)))
score = f1_score(cm_correct_labels, cm_predictions, labels=range(len(CLASSES)), average='macro')
precision = precision_score(cm_correct_labels, cm_predictions, labels=range(len(CLASSES)), average='macro')
recall = recall_score(cm_correct_labels, cm_predictions, labels=range(len(CLASSES)), average='macro')
cmat = (cmat.T / cmat.sum(axis=1)).T # normalized
display_confusion_matrix(cmat, score, precision, recall)
print('f1 score: {:.3f}, precision: {:.3f}, recall: {:.3f}'.format(score, precision, recall))
```

/usr/local/lib/python3.8/site-packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))





f1 score: 0.827, precision: 0.866, recall: 0.807

Predictions

```
In [17]: test_ds = get_test_dataset(ordered=True) # since we are splitting the dataset and iterating separately on images and ids

print('Computing predictions...')
test_images_ds = test_ds.map(lambda image, idnum: image)
probabilities = model.predict(test_images_ds, steps=TEST_STEPS)
predictions = np.argmax(probabilities, axis=-1)
print(predictions)

print('Generating submission.csv file...')
test_ids_ds = test_ds.map(lambda image, idnum: idnum).unbatch()
test_ids = next(iter(test_ids_ds.batch(NUM_TEST_IMAGES))).numpy().astype('U') # all in one batch
np.savetxt('submission.csv', np.rec.fromarrays([test_ids, predictions]), fmt=['%s', '%d'], delimiter=',', header='id,label')
!head submission.csv
```

```
Computing predictions...
58/58 [=====] - 15s 258ms/step
[ 14  83 103 ...  49  47 103]
Generating submission.csv file...
id,label
0b9afbd2,14
c37a6f3e9,83
00e4f514e,103
1c4736dea,28
252d840db,67
dfc9c6a23,103
53cf6586,48
541c4d41e,102
59d1b6146,70
```

Visual validation

```
In [18]: dataset = get_validation_dataset()
dataset = dataset.unbatch().batch(20)
batch = iter(dataset)
```

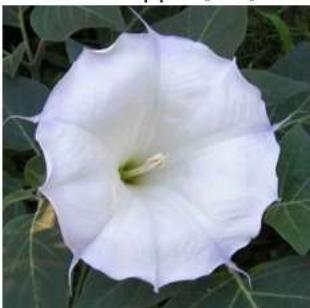
```
In [19]: # run this cell again for next set of images
images, labels = next(batch)
probabilities = model.predict(tf.cast(images, tf.float32))
predictions = np.argmax(probabilities, axis=-1)
display_batch_of_images((images, labels), predictions)
```

2023-10-22 21:42:18.746861: E tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] model_pruner failed: INVALID_ARGUMENT: Graph does not contain terminal node AssignAddVariableOp.

2023-10-22 21:42:18.939716: E tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] model_pruner failed: INVALID_ARGUMENT: Graph does not contain terminal node AssignAddVariableOp.

1/1 [=====] - 9s 9s/step

thorn apple [OK]



wild pansy [OK]



tree poppy [OK]



king protea [OK]



wild rose [OK]



pink primrose [OK]



iris [OK]



common dandelion [OK]



daisy [OK]



yellow iris [OK]



wild geranium [OK]



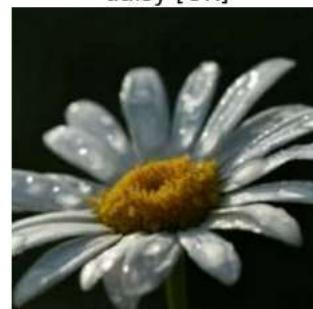
sword lily [OK]



hibiscus [OK]



daisy [OK]



common dandelion [OK]



orange dahlia [OK]



common dandelion [OK]



sunflower [NO→black-eyed susan]



petunia [OK]



cyclamen [NO→azalea]

